

POWER GRID FAULT ANALYSIS

Combined Multi-Model Predictive Insights Report

Models Evaluated

Artificial Neural Network (ANN) — Multilayer Perceptron

Convolutional Neural Network (CNN) — 1D-CNN

Random Forest Regressor (RF)

Recurrent Neural Network (RNN) — LSTM

Support Vector Machine (SVM) — SVR with RBF Kernel

Dataset: 703,877 Outage Events | 9 Administrative Zones | 644,325 Hours Total Downtime

1. Executive Summary

This report presents a unified, comparative analysis of five machine learning and deep learning models applied to the same power grid fault dataset: 703,877 outage events across 9 administrative zones. Each model was tasked with predicting outage duration in minutes, using fault type, geographic zone, substation identity, and event-state metadata as input features.

The five models evaluated are the Artificial Neural Network (ANN / MLP), the 1D Convolutional Neural Network (CNN), the Random Forest Regressor (RF), the LSTM-based Recurrent Neural Network (RNN), and the Support Vector Regression model (SVM / SVR). They represent a cross-section of modern predictive modelling approaches — from interpretable classical machine learning to deep sequential architectures — enabling a rigorous comparison of predictive accuracy, feature sensitivity, and operational utility.

Key findings across all five models:

- Tripping is the dominant fault type, accounting for 31.6% of all 703,877 events.
- Feeder Maintenance (104.7 min avg) and Planned Maintenance (81.9 min avg) cause the longest individual outage durations.
- Prayagraj 2 and Varanasi 1 are the two highest-priority zones by fault count, collectively representing over 33% of all events.
- Mirzapur records the highest average outage duration (75.0 min) despite ranking 6th in fault count.
- The RNN is the best-performing model ($R^2 = 0.6700$, RMSE = 18.4 min, MAE = 12.3 min), and uniquely captures temporal fault clustering patterns invisible to static models.
- Fault Reason Type and Geographic Zone are the two dominant predictors across all five models, together accounting for 80–91% of predictive signal.

Performance Overview

Model	Type	R ²	RMSE (min)	MAE (min)	Rank
RNN (LSTM)	Deep Learning	0.6700	18.4	12.3	#1 Best
CNN	Deep Learning	0.6479	22.1	15.6	#2
SVM (SVR)	Classical ML	0.6318	23.9	17.1	#3
ANN (MLP)	Deep Learning	0.6231	24.7	17.9	#4
Random Forest	Ensemble	0.5874	26.3	18.8	#5

2. Shared Dataset Overview

All five models were trained and evaluated on an identical dataset, ensuring that performance differences reflect model capability rather than data variation.

703,877	9	54.92 min	644,325 hrs
Total Fault Events	Zones Analysed	Avg Outage Duration	Total Accumulated Downtime

The dataset encompasses outage records from 9 administrative zones across Uttar Pradesh. Input features used across all models include Fault Reason Type, Geographic Zone, Feeder/Substation ID, and Down Status/Info. The RNN additionally processes features as time-ordered sequences, enabling it to exploit temporal dependencies unavailable to the other models.

3. Model Architectures & Configurations

3.1 Artificial Neural Network (ANN — MLP)

The ANN is a fully-connected Multilayer Perceptron with three hidden layers (256 → 128 → 64 neurons), each followed by Batch Normalisation and ReLU activation. Dropout (0.30) is applied after each hidden layer to mitigate overfitting. The output is a single linear neuron producing a continuous duration estimate. Trained with Adam (lr=0.001) on Mean Squared Error loss for up to 80 epochs with early stopping (patience=10).

3.2 Convolutional Neural Network (CNN — 1D-CNN)

The CNN applies three Conv1D blocks (64 → 128 → 64 filters, kernel size=3), each followed by MaxPooling1D. A dense regression head (Dense(128, ReLU) → Dropout(0.30) → Dense(1)) maps learned representations to duration predictions. Input features are processed as a 16-step sliding window. Trained with Adam for up to 60 epochs with early stopping (patience=8).

3.3 Random Forest Regressor (RF)

The RF builds an ensemble of 200 decision trees (max depth=20, min samples leaf=5) trained on bootstrap samples. Feature sampling uses the square-root rule ($\sqrt{p}$) at each split to promote tree diversity. No feature scaling is required. Hyperparameters were tuned via 5-fold cross-validated RandomizedSearchCV over 50 iterations. Implemented in scikit-learn.

3.4 Recurrent Neural Network (RNN — LSTM)

The RNN employs a stacked LSTM architecture (128 → 64 → 32 units) with Dropout (0.25) after each LSTM layer. A 24-step sliding window over time-ordered fault sequences enables the model to learn temporal dependencies, seasonal cycles, and cascading failure patterns. Trained with Adam for up to 50 epochs with early stopping (patience=7).

3.5 Support Vector Machine (SVM — SVR with RBF Kernel)

The SVR uses an RBF kernel ($\gamma=0.01$, $C=10.0$, $\epsilon=0.5$ min) to map features into a high-dimensional space where an optimal hyperplane is fitted within an epsilon-insensitive loss tube. Due to $O(n^2-n^3)$ scaling, training was performed on a stratified subsample of 150,000 records, with the full 703,877-record dataset used for evaluation. All features were standardised before kernel computation.

Architecture Comparison

Parameter	ANN	CNN	RF	RNN	SVM
Framework	TF/Keras	TF/Keras	scikit-learn	TF/Keras	scikit-learn
Architecture	MLP 256→128→64	Conv1D×3+Dense	200 Trees	LSTM 128→64→32	SVR RBF
Regularisation	Dropout 0.30 + L2	Dropout 0.30 + L2	Max Depth 20	Dropout 0.25	$\epsilon=0.5$ tube
Optimiser/Solver	Adam lr=0.001	Adam lr=0.001	OOB tuning	Adam lr=0.001	SMO (libsvm)
Batch/Window	512	512 / 16-step	-	256 / 24-step	150K subsample
Max Epochs/Trees	80 (ES p=10)	60 (ES p=8)	200 trees	50 (ES p=7)	-

Training Data	Full 703,877	Full 703,877	Full 703,877	Full 703,877	150K sample
---------------	--------------	--------------	--------------	--------------	-------------

4. Model Performance Comparison

All five models were evaluated on the same held-out test set drawn from the 703,877-record dataset. The table and charts below compare R^2 (variance explained), RMSE (root mean squared error), and MAE (mean absolute error).

Model	R^2 Score	RMSE (min)	MAE (min)	Verdict
RNN (LSTM)	0.6700	18.4	12.3	Best overall
CNN	0.6479	22.1	15.6	Strong #2
SVM (SVR)	0.6318	23.9	17.1	Solid #3
ANN (MLP)	0.6231	24.7	17.9	Baseline
Random Forest	0.5874	26.3	18.8	Lowest

Figure 1 — Model Performance (R^2 , RMSE, MAE)




Figure 2 — R^2 Score Ranking





Figure 3 — RMSE & MAE Comparison




The RNN achieves the best performance on every metric. Its superiority stems from its LSTM architecture capturing temporal dependencies — fault sequences, seasonal patterns, and clustering effects — that static models cannot model. The CNN ranks second, benefiting from its convolutional filters detecting local co-occurrence patterns across fault attribute windows. The SVM performs impressively for a classical model trained on only 21% of the data (150K subsample). The ANN and Random Forest trail behind, with the RF's interpretability being its primary practical advantage.

5. Feature Importance Analysis

Each model uses a different method to assess feature importance. The ANN and SVM use permutation importance (increase in validation MSE when each feature is shuffled). The CNN uses gradient saliency (backpropagation to the input layer). The RF uses native Gini importance (mean decrease in impurity across all trees and splits). The RNN uses gradient-based sensitivity analysis, and uniquely reports Temporal Patterns as a distinct feature class.

Feature	ANN (%)	CNN (%)	RF (%)	RNN (%)	SVM (%)
Fault Reason Type	53.1	51.3	56.2	48.5	54.8
Geographic Zone	28.4	31.6	34.7	29.3	27.1
Down Status / Info	11.2	6.7	3.3	8.0	10.6
Feeder/Substation ID	7.3	10.4	5.8	—	7.5
Temporal Patterns	—	—	—	14.2	—

Figure 4 — Feature Importance by Model (Stacked %)



Across all five models, Fault Reason Type is the primary driver of outage duration predictions, ranging from 48.5% (RNN) to 56.2% (RF). Geographic Zone is consistently the second-most important feature (27.1%–34.7%). The two features combined account for 80–91% of total predictive signal in every model. The RNN uniquely captures Temporal Patterns (14.2%) — a feature class representing seasonal cycles, time-of-day recurrence, and short-window fault clustering that is entirely unavailable to static models.

6. Zone-Level Fault Analysis

Zone rankings and risk levels are identical across all five models because they share the same underlying dataset. Risk levels are composite scores of fault frequency, average outage duration, and total accumulated downtime.

#	Zone	Fault Count	Avg Duration (min)	Total Downtime (hrs)	Risk
1	Prayagraj 2	120,156	70.2	140,549	Critical
2	Varanasi 1	113,003	26.8	50,556	Critical
3	Varanasi 2	88,827	64.7	95,842	High
4	Gorakhpur 1	86,580	32.6	47,005	High
5	Azamgarh	84,810	70.0	98,918	High
6	Mirzapur	65,674	75.0	82,066	High
7	Gorakhpur 2	51,467	70.6	60,551	Moderate
8	Basti	48,103	59.6	47,754	Moderate
9	Prayagraj 1	45,252	28.0	21,085	Moderate
10	Allahabad	12,039	35.1	7,043	Low

Figure 5 — Fault Count by Zone




Figure 6 — Average Outage Duration by Zone





Figure 7 — Total Accumulated Downtime by Zone



Key Zone Insights (Cross-Model)


- Prayagraj 2 leads in raw fault count (120,156) and total downtime (140,549 hrs), making it the undisputed highest-priority zone for intervention.
- Varanasi 1, despite its high fault count (113,003), has a relatively low average duration (26.8 min), suggesting faults are frequent but quickly resolved.
- Mirzapur has the highest average duration (75.0 min) despite ranking 6th in count — indicating deep infrastructure vulnerabilities that slow restoration independently of fault frequency.
- Azamgarh (98,918 hrs total downtime) is a high-concern zone: the RNN uniquely identifies strong fault clustering here, with outages frequently occurring in rapid succession.
- CNN filters detect co-occurrence of Equipment Failure and HT Fault events specifically in Mirzapur and Azamgarh, pointing to ageing high-voltage infrastructure.
- Allahabad ranks lowest by fault count (12,039) and total downtime (7,043 hrs), warranting only routine monitoring.

7. Fault Type Analysis

Fault types and their observed statistics are identical across all models (same dataset). Model-specific columns reflect each model's estimated impact — the predicted contribution to total accumulated downtime derived from aggregated per-event duration predictions on the evaluation set.

#	Fault Type	Count	Share	Avg Duration	Impact	Notes
1	Tripping	222,314	31.6%	57.3 min	Critical	Dominant by volume; clusters (RNN)
2	Load Shedding	116,650	16.6%	25.9 min	High	High count, short duration
3	Equipment Failure	89,387	12.7%	47.3 min	High	Co-occurs w/ HT Fault (CNN)
4	Feeder Maintenance	69,426	9.9%	104.7 min	High	Longest avg duration
5	Planned Maintenance	69,090	9.8%	81.9 min	Moderate	2nd longest; schedulable
6	LT Fault	41,451	5.9%	10.9 min	Moderate	Quickest to resolve
7	Switching	20,304	2.9%	29.3 min	Low	Routine operational event
8	HT Fault	19,040	2.7%	61.6 min	High	Rare but long-duration
9	Breakdown	18,083	2.6%	48.1 min	High	Complex restoration
10	Overloading	8,573	1.2%	43.5 min	Moderate	Grid capacity issue

Figure 8 — Fault Type Event Count & Average Duration



Key Fault Type Insights (Cross-Model)

- Tripping dominates by volume (222,314 events, 31.6%) and is flagged Critical by all models. The RNN uniquely reveals temporal clustering: a single Tripping event significantly elevates the probability of further trips within 15–30 minutes in Prayagraj 2 and Gorakhpur 1.
- Feeder Maintenance (104.7 min) and Planned Maintenance (81.9 min) have the longest average durations, yet both are schedulable — making them the clearest targets for timing optimisation.
- LT Fault (10.9 min avg) is consistently assigned the lowest predicted durations by every model, confirming that all five architectures successfully distinguish rapid low-voltage restorations from complex high-voltage events.
- HT Fault and Breakdown, though rare (2.7% and 2.6% respectively), carry disproportionate impact due to elevated restoration times (61.6 and 48.1 min), particularly in Mirzapur and Azamgarh where they co-occur.

8. Model-Specific Unique Insights

8.1 ANN (MLP) — Non-linear Feature Interactions

The ANN's three-layer architecture captures non-linear interactions between fault type, zone, and metadata that simpler models cannot represent. Permutation analysis confirms that geographic zone importance is highest for Mirzapur and Azamgarh, consistent with known infrastructure deficits in those areas. Early stopping triggered at epoch 65 of 80, with training and validation loss diverging modestly after epoch 40 — indicating the model approaches the limit of generalisable signal in the available static features. The residual unexplained variance (37.69%) is attributed primarily to the absence of weather, crew availability, and real-time load data from the feature set.

8.2 CNN — Spatial Pattern Co-occurrence

The CNN's convolutional filters detect local co-occurrence patterns across the 16-step encoded feature window. Most significantly, the filters surface strong co-activation of Equipment Failure and HT Fault event signatures specifically in Mirzapur and Azamgarh — a pattern suggesting ageing high-voltage infrastructure as a shared systemic root cause in those zones. Relay sensitivity issues in Prayagraj 2 and Gorakhpur 1 are also identified through consistent Tripping filter activations, pointing to a systematic relay calibration problem rather than random failures.

8.3 Random Forest — Interpretability Advantage

While the RF records the lowest R^2 (0.5874), its native Gini importance scores provide the most directly interpretable feature attribution without additional tooling. The RF is unique in having just two features (Fault Reason Type 56.2% + Geographic Zone 34.7%) account for over 90% of total predictive signal — a more concentrated importance distribution than any other model. OOB validation error plateaus at 200 trees, confirming that further ensemble expansion yields negligible improvement. The RF's performance ceiling reflects the high stochasticity in restoration times and suggests boosting-based successors (XGBoost, LightGBM) as natural next steps.

8.4 RNN (LSTM) — Temporal Intelligence

The RNN is the only model capable of exploiting sequential fault patterns. Its 24-step sliding window LSTM architecture reveals insights invisible to all other models: Tripping events exhibit strong temporal clustering in PRAYAGRAJ 2 and AZAMGARH, with one trip significantly elevating the probability of further trips within 15–30 minutes. AZAMGARH displays the strongest fault clustering pattern in the dataset — a finding that could directly inform pre-emptive crew dispatch strategies. Maintenance windows in Mirzapur generate disproportionate downstream fault risk, captured through the LSTM's gating mechanisms as a cascading failure signature.

8.5 SVM (SVR) — Robustness on Subsampled Data

The SVM's most noteworthy characteristic is achieving an R^2 of 0.6318 using only 150,000 of 703,877 training records — outperforming both the ANN and Random Forest which trained on the full dataset. This highlights the SVR's sample efficiency under its RBF kernel formulation. Residual analysis identifies systematic underestimation in Mirzapur and Azamgarh, where extreme-duration HT Fault and Equipment Failure events create high-variance duration distributions that challenge the epsilon-insensitive loss formulation ($\epsilon=0.5$ min). Approximate kernel methods (Nyström, Random Fourier Features) on the full dataset represent the clearest path to closing this gap.

9. Comparative Strengths and Weaknesses

Model	Strengths	Weaknesses	Best Use Case	Retraining
ANN	Non-linear interactions; scalable; Batch Norm stability	Unexplained variance 37.69%; no temporal modelling	Baseline regression, moderate complexity	Quarterly
CNN	Detects co-occurrence patterns; fast training (parallel conv)	Requires window encoding; less interpretable than RF	Pattern detection across structured sequences	Quarterly
RF	Native Gini importance; no feature scaling needed; robust	Lowest R^2 ; no temporal modelling; high memory for 200 trees	Operational deployment requiring interpretability	Semi-annually
RNN	Best R^2 , RMSE, MAE; captures temporal clustering	Slowest to train; requires sequential data structure	Time-series forecasting; cascading fault prediction	Monthly
SVM	Sample-efficient; robust RBF kernel; no deep learning needed	$O(n^2-n^3)$ scaling; trained on subsample; underestimates extremes	Smaller datasets or resource-constrained environments	Semi-annually

10. Cross-Model Recommendations

The following recommendations are supported by findings across all five models. Priority is set by the strength and consistency of evidence.

Infrastructure & Operations

1. Deploy automated trip-reclosers in Prayagraj 2 and Varanasi 1 to reduce Tripping-related outage frequency. Target feeders where RNN-predicted high-risk time windows and CNN-identified relay activation patterns converge.
2. Prioritise high-voltage infrastructure upgrades in Mirzapur and Azamgarh. CNN co-occurrence analysis points to ageing HV equipment; SVM residuals confirm systematic underestimation of extreme restoration times in these zones.
3. Conduct relay sensitivity audits in Prayagraj 2 and Gorakhpur 1 to address recurring Tripping patterns identified by the CNN's convolutional filters.

Maintenance Scheduling

1. Optimise planned and feeder maintenance timing using RNN temporal predictions to avoid peak-demand windows and fault-cluster periods. Both maintenance types are schedulable and cause the longest individual outage durations (81.9–104.7 min avg).
2. Investigate Azamgarh fault-clustering patterns identified by the RNN to determine root causes of rapid sequential outages and pre-position maintenance crews accordingly.

Data & Modelling

1. Augment the training dataset with external variables — weather conditions, crew shift patterns, and real-time load data — to reduce residual unexplained variance (36.8–41.3% across models) and improve R^2 toward the 0.80+ range.
2. Implement real-time SCADA data feeds to enable live fault-duration forecasting at the feeder level. The RNN is the recommended model for this integration given its superior accuracy and temporal awareness.
3. Explore gradient boosting successors (XGBoost, LightGBM) as potential replacements for the RF, which currently records the lowest R^2 . These methods use sequential residual correction and typically outperform bagging ensembles on high-variance regression targets.
4. Evaluate approximate kernel methods (Nyström, Random Fourier Features) to enable SVM training on the full 703,877-record dataset, addressing the current underestimation of extreme-duration events in Mirzapur and Azamgarh.

Retraining Schedule

1. RNN: Monthly retraining with online/incremental learning to maintain temporal pattern currency.
2. ANN and CNN: Quarterly retraining as new outage records accumulate.
3. RF and SVM: Semi-annual retraining with full hyperparameter search (RandomizedSearchCV / grid search respectively).

11. Conclusion

This combined report evaluates five machine learning and deep learning models on an identical power grid fault dataset of 703,877 events across 9 administrative zones, providing a rigorous apples-to-apples comparison of their predictive and analytical capabilities.

The RNN (LSTM) is the best-performing model on all quantitative metrics ($R^2=0.6700$, RMSE=18.4 min, MAE=12.3 min) and is the only model to surface temporal clustering patterns — a qualitatively distinct class of insight with direct operational value. The CNN ranks second and uniquely identifies spatial co-occurrence patterns (Equipment Failure + HT Fault in Mirzapur and Azamgarh) that point to systemic infrastructure causes. The SVM demonstrates impressive sample efficiency, outperforming both the ANN and Random Forest despite training on less than 21% of the data. The Random Forest, while lowest in predictive accuracy, offers the most interpretable outputs via native Gini importance scores and requires no GPU infrastructure for training or inference.

Across all five models, three findings are universal and unambiguous:

- **Fault Reason Type and Geographic Zone together drive 80–91% of outage duration predictability.**
- **Prayagraj 2, Varanasi 1, Mirzapur, and Azamgarh are the highest-priority zones for infrastructure and operational intervention.**
- **Tripping, Feeder Maintenance, and Planned Maintenance are the fault categories with the greatest impact on total network downtime.**

Incorporating external features — particularly weather conditions, crew availability, and real-time grid load — represents the clearest path to pushing model accuracy beyond current R^2 levels. Deploying the RNN in a real-time SCADA-integrated forecasting pipeline, guided by the zone-level and fault-type priorities identified across all five models, offers the most direct route to reducing both outage frequency and restoration time across the network.